

Assignment 2

Data cleaning

Load data and find feature present in both approved and rejected data.

```
In [1]: import pandas as pd

approved_csv = pd.read_csv("LoanStats3a.csv", skiprows=1)
rejected_csv = pd.read_csv("RejectStatsA.csv", skiprows=1)

### Mapping ###
# 'loan_amnt' - 'Amount Requested'
# 'issue_d' - 'Application Date'
# 'purpose', 'title' - 'Loan Title'
# '' - 'Risk_Score' (no match)
# 'dti' - 'Debt-To-Income Ratio'
# 'zip_code' - 'Zip Code'
# 'addr_state' - 'State'
# 'policy_code' - 'Policy Code'
# 'emp_length' - 'Employment Length'
approved_raw_data = approved_csv[['loan_amnt', 'issue_d', 'purpose', 'dti', 'zip_code',
                                     'addr_state', 'policy_code', 'emp_length']]
rejected_raw_data = rejected_csv[['Amount Requested', 'Application Date', 'Loan Title',
                                   'Debt-To-Income Ratio', 'Zip Code', 'State',
                                   'Policy Code', 'Employment Length']]
rejected_raw_data.columns = ['loan_amnt', 'issue_d', 'purpose', 'dti', 'zip_code',
                              'addr_state', 'policy_code', 'emp_length']

/Library/Frameworks/Python.framework/Versions/3.6/lib/python3.6/site-packages/IPython/core/interactiveshell.py:2785: DtypeWarning: Columns (0, 47) have mixed types. Specify dtype option on import or set low_memory=False.
interactivity=interactivity, compiler=compiler, result=result)
```

Count the number of entries available in each set, and display a small sample.

```
In [2]: print ("The dataset contain {:d} approved data.".format(approved_raw_data.shape[0]))
        print ("The dataset contain {:d} rejected data.".format(rejected_raw_data.shape[0]))

        print ("\n")
        print ("Sample from approved data:")
        print (approved_raw_data.sample(5))
        print ("\n")
        print ("Sample from rejected data:")
        print (rejected_raw_data.sample(5))
```

The dataset contain 42538 approved data.
 The dataset contain 755491 rejected data.

Sample from approved data:

te \	loan_amnt	issue_d	purpose	dti	zip_code	addr_sta
31404	6250.0	Mar-2010	debt_consolidation	18.04	342xx	
FL						
25328	5000.0	Sep-2010	debt_consolidation	16.78	010xx	
MA						
11346	3050.0	Jul-2011	debt_consolidation	16.65	799xx	
TX						
21971	11500.0	Dec-2010	debt_consolidation	15.32	923xx	
CA						
32074	8000.0	Feb-2010	small_business	1.65	770xx	
TX						

	policy_code	emp_length
31404	1.0	10+ years
25328	1.0	8 years
11346	1.0	< 1 year
21971	1.0	5 years
32074	1.0	1 year

Sample from rejected data:

_state \	loan_amnt	issue_d	purpose	dti	zip_code	addr
3670	25000.0	2007-12-02	2Elisaiah	34.2%	300xx	
GA						
748504	25000.0	2012-12-26	credit_card	10.18%	761xx	
TX						
557466	25000.0	2012-07-07	debt_consolidation	19.41%	785xx	
TX						
336713	3000.0	2011-09-03	car	51.8%	352xx	
AL						
83793	10000.0	2009-12-01		0%	412xx	
KY						

	policy_code	emp_length
3670	0	5 years
748504	0	< 1 year
557466	0	< 1 year
336713	0	< 1 year
83793	0	10+ years

Count how may entries have missing value for both set.

```
In [3]: nan_count = pd.DataFrame({"approved": approved_raw_data.isnull().sum(),
                                "approved %": approved_raw_data.isnull().sum()
                                / len(approved_raw_data),
                                "rejected": rejected_raw_data.isnull().sum(),
                                "rejected %": rejected_raw_data.isnull().sum()
                                / len(rejected_raw_data)})
print (nan_count)
```

	approved	approved %	rejected	rejected %
loan_amnt	3	0.000071	0	0.000000
issue_d	3	0.000071	0	0.000000
purpose	3	0.000071	14	0.000019
dti	3	0.000071	0	0.000000
zip_code	3	0.000071	22	0.000029
addr_state	3	0.000071	21	0.000028
policy_code	3	0.000071	0	0.000000
emp_length	1115	0.026212	8130	0.010761

In general, the dataset is relatively complet, with most of the missing value occur in the 'emp_length' feature. However, it is still less than 3% of missing value for the approved set and about 1% for the rejected set. This means it is relatively safe to simply drop the entries with NaN value with low possibility of introducing a significant bias.

```
In [4]: approved_raw_data = approved_raw_data.dropna()
rejected_raw_data = rejected_raw_data.dropna()
```

Data processing

Feature 'issue_d' and 'dti' are coded in different format in the approved and rejected set. We will need to standardize the format.

```
In [5]: import copy

# deep copy raw data into intermediate processing set 0
approved_data_0 = copy.deepcopy(approved_raw_data)
rejected_data_0 = copy.deepcopy(rejected_raw_data)
```

```
In [6]: from datetime import datetime
import numpy as np

def date_string_to_datetime(row):
    try:    # for approved dataset
        return datetime.strptime(row['issue_d'], "%b-%Y")
    except: # for rejected dataset
        return datetime.strptime(row['issue_d'], "%Y-%m-%d")

def extract_year_lambda(row):
    return date_string_to_datetime(row).year

def extract_month_lambda(row):
    return date_string_to_datetime(row).month

def dti_percent_to_decimal(row):
    return np.float64(row['dti'].strip("%"))

approved_data_0['issue_y'] = approved_data_0.apply(extract_year_lambda,
axis=1)
approved_data_0['issue_m'] = approved_data_0.apply(extract_month_lambda,
axis=1)
approved_data_0 = approved_data_0.drop(columns=['issue_d'])

rejected_data_0['issue_y'] = rejected_data_0.apply(extract_year_lambda,
axis=1)
rejected_data_0['issue_m'] = rejected_data_0.apply(extract_month_lambda,
axis=1)
rejected_data_0['dti'] = rejected_data_0.apply(dti_percent_to_decimal, a
xis=1)
rejected_data_0 = rejected_data_0.drop(columns=['issue_d'])
```

```
In [7]: print ("Sample from approved data:")
print (approved_data_0.sample(5))
print ("\n")
print ("Sample from rejected data:")
print (rejected_data_0.sample(5))
```

Sample from approved data:

	loan_amnt	purpose	dti	zip_code	addr_state	policy
34453	25000.0	wedding	9.96	070xx	NJ	
1.0						
39199	15000.0	debt_consolidation	15.10	275xx	NC	
1.0						
40935	15000.0	debt_consolidation	20.44	334xx	FL	
1.0						
28389	14000.0	debt_consolidation	20.98	750xx	TX	
1.0						
17704	2400.0	major_purchase	18.47	749xx	OK	
1.0						

	emp_length	issue_y	issue_m
34453	4 years	2009	11
39199	< 1 year	2008	3
40935	8 years	2009	10
28389	2 years	2010	7
17704	10+ years	2011	3

Sample from rejected data:

	loan_amnt	purpose	dti	zip_code	addr_state
414756	3000.0	other	6.79	583xx	MN
209994	25000.0	small_business	2.67	786xx	TX
390181	3000.0	car	0.00	300xx	GA
503258	25000.0	Debt Consolidation Loan	14.76	303xx	GA
156202	25000.0	debt_consolidation	83.47	127xx	NY

	policy_code	emp_length	issue_y	issue_m
414756	0	< 1 year	2011	12
209994	0	< 1 year	2011	1
390181	0	< 1 year	2011	11
503258	0	9 years	2012	5
156202	0	< 1 year	2010	9

Feature 'purpose' is a string, containing discription of the intened use of the loan. We need to find a way to represent that data in a machine-friendly way.

```
In [8]: # in the approved set, the purpose data is well organized label
print ('purpose' feature count in approved dataset:)
print (approved_data_0[['purpose', 'policy_code']].groupby('purpose').count())
print ("\n")

# in the rejected set, the purpose data is messy text string
print ("Sample of 'purpose' feature count in rejected dataset:")
print (rejected_data_0[['purpose', 'policy_code']].groupby('purpose').count().sample(10))
```

'purpose' feature count in approved dataset:

	policy_code
purpose	
car	1563
credit_card	5344
debt_consolidation	19363
educational	413
home_improvement	3099
house	412
major_purchase	2238
medical	726
moving	603
other	4259
renewable_energy	98
small_business	1946
vacation	368
wedding	991

Sample of 'purpose' feature count in rejected dataset:

	policy_code
purpose	
Tradition pay-off	1
Business Supplies	1
Loan for student loan and paying bills	1
BE ABLE TO BE BACK ON TRACK.....	1
For College	2
terrylou	1
Cadillac 96	1
Need Quick-In Escrow	1
Replace Privacy Fence and Update Kitchen	1
looking to expand	1

To clean the 'purpose' feature data in the rejected set, we will use KNN to cluster the text strings in rejected set to the labels in approved set.

```

In [9]: from sklearn.feature_extraction.text import TfidfVectorizer
        from sklearn.cluster import KMeans
        from sklearn.metrics import adjusted_rand_score
        import numpy as np

np.random.seed(5)

# extract 'purpose' feature data for both sets
rejected_purpose_text = list(rejected_data_0['purpose'])
approved_purpose_label = list(approved_data_0.groupby('purpose').groups.
keys())

# build knn cluster model using all text in the rejected set
vectorizer = TfidfVectorizer(stop_words='english')
X = vectorizer.fit_transform(rejected_purpose_text)
true_k = 14
knn_cluster_model = KMeans(n_clusters=true_k, init='k-means++', max_iter
=100, n_init=1)
knn_cluster_model.fit(X)

# print a summary of the clusters
print ("Top terms in each cluster:")
order_centroids = knn_cluster_model.cluster_centers_.argsort()[:, :-1]
terms = vectorizer.get_feature_names()
for i in range (true_k):
    print ("Cluster {:d}: {:s}, {:s}, {:s}".format(* [i] + [terms[ind] f
or ind in order_centroids[i, :3]] ))

# map the approved set label to the cluster model using predict method
Y = vectorizer.transform(approved_purpose_label)
prediction = knn_cluster_model.predict(Y)
label_mapping = pd.DataFrame({'original_label': approved_purpose_label,
                              'predicted_cluster': prediction,
                              'predicted_cluster_top_label': [terms[orde
r_centroids[c][0]] for c in prediction]})
print ("\n")
print ("Predicted mapping between approved set 'purpose' label and rejec
ted set 'purpose' text text")
print (label_mapping)
unmatched_cluster = list(set(range(14)) - set(prediction))
print ("Unmatched clusters:", unmatched_cluster)
print ("With the top label in cluster being:", [terms[order_centroids[c]
[0]] for c in unmatched_cluster])

```


Top terms in each cluster:

Cluster 0: moving, loan, expenses
 Cluster 1: debt_consolidation, factoring, facelift
 Cluster 2: home_improvement, loan, zzzzzgirl
 Cluster 3: credit_card, loan, zzzzzgirl
 Cluster 4: loan, consolidation, debt
 Cluster 5: major_purchase, loan, zzzzzgirl
 Cluster 6: medical, expenses, loan
 Cluster 7: car, loan, need
 Cluster 8: small_business, loan, zzzzzgirl
 Cluster 9: house, loan, buy
 Cluster 10: wedding, loan, expenses
 Cluster 11: vacation, loan, dream
 Cluster 12: consolidate, debt, credit
 Cluster 13: bills, pay, medical

Predicted mapping between approved set 'purpose' label and rejected set 'purpose' text text

	original_label	predicted_cluster	predicted_cluster_top_label
0	car	7	car
1	credit_card	3	credit_card
2	debt_consolidation	1	debt_consolidation
3	educational	4	loan
4	home_improvement	2	home_improvement
5	house	9	house
6	major_purchase	5	major_purchase
7	medical	6	medical
8	moving	0	moving
9	other	4	loan
10	renewable_energy	4	loan
11	small_business	8	small_business
12	vacation	11	vacation
13	wedding	10	wedding

Unmatched clusters: [12, 13]

With the top label in cluster being: ['consolidate', 'bills']

By applying KNN clustering on the rejected set, we are able to map most of the label in the approved set to the cluster of the rejected set. The exceptions we have are for the 'educational', 'other', and 'renewable_energy' label. However, as 'educational' and 'renewable_energy' are relatively small group (413 and 98 entries compare to the over 40k dataset) and 'other' simply means other, we can group these labels together as 'other'. For the unmatched cluster 12 and 13, its top keyword being 'consolidate' and 'bills' also fit to the description of 'other'.

```
In [10]: import copy

# deep copy intermediate processing set 0 into intermediate processing set 1
approved_data_1 = copy.deepcopy(approved_data_0)
rejected_data_1 = copy.deepcopy(rejected_data_0)
```

```
In [11]: import swifter

def text_to_cluster_code(row):
    text = vectorizer.transform([row['purpose']])
    cluster_code = knn_cluster_model.predict(text)[0]
    #cluster_code = cluster_code[0]
    if cluster_code is 12 or cluster_code is 13:
        cluster_code = 4
    return cluster_code

approved_data_1['purpose'] = approved_data_1.swifter.apply(text_to_cluster_code, axis=1)
rejected_data_1['purpose'] = rejected_data_1.swifter.apply(text_to_cluster_code, axis=1)
```

```
Pandas Apply: 100%|██████████| 41423/41423 [00:22<00:00, 1829.03it/s]
Pandas Apply: 100%|██████████| 747325/747325 [07:31<00:00, 1654.33it/s]
```

```
In [12]: print ("Sample from approved data:")
print (approved_data_1.sample(5))
print ("\n")
print ("Sample from rejected data:")
print (rejected_data_1.sample(5))
```

Sample from approved data:

	loan_amnt	purpose	dti	zip_code	addr_state	policy_code	emp_l
length \							
37835	5650.0	1	6.18	927xx	CA	1.0	10+
years							
31779	5000.0	1	8.93	107xx	NY	1.0	10+
years							
23927	20000.0	2	11.09	189xx	PA	1.0	5
years							
3904	6000.0	11	10.98	606xx	IL	1.0	1
year							
17841	6000.0	8	13.06	487xx	MI	1.0	1
year							

	issue_y	issue_m
37835	2009	1
31779	2010	3
23927	2010	11
3904	2011	11
17841	2011	4

Sample from rejected data:

	loan_amnt	purpose	dti	zip_code	addr_state	policy_code	emp_
length \							
518432	15000.0	1	14.00	280xx	NC	0	8
years							
141699	30000.0	1	39.45	154xx	PA	0	10+
years							
495626	25000.0	5	5.78	212xx	MD	0	<
1 year							
629279	5000.0	1	7.87	549xx	WI	0	<
1 year							
737255	25000.0	4	67.79	553xx	MN	0	<
1 year							

	issue_y	issue_m
518432	2012	5
141699	2010	7
495626	2012	4
629279	2012	9
737255	2012	12

Feature 'emp_length' now is a categorical data, however, as the nature of the data is numerical, we will convert it into numerical data. Feature 'purpose', 'addr_state', 'issue_y', and 'issue_m' are categorical data, and we will represent them using dummy variables. Feature 'zip_code' is between a categorical and numerical data, and is very difficult to use directly as a feature without more pre-processing. We will drop it for simplicity purpose. In addition, we will normalize the data

```
In [35]: import copy

# deep copy intermediate processing set 1 into intermediate processing set 2
approved_data_2 = copy.deepcopy(approved_data_1)
rejected_data_2 = copy.deepcopy(rejected_data_1)

# at this point, the remaining job is only to clean up, we can join the two sets
# policy_code=1 represent approved, policy_code=0 represent rejected
data = pd.concat([approved_data_2, rejected_data_2])
```

```
In [36]: from sklearn import preprocessing

# Convert originally categorical data employment length to numerical data
data['emp_length'].replace(to_replace='[^0-9]+', value='', inplace=True,
                           regex=True)
data['emp_length'] = data['emp_length'].astype(int)

# generate dummy variables for 'purpose', 'addr_state', 'issue_y', and 'issue_m'
data = pd.get_dummies(data, columns=['purpose', 'addr_state', 'issue_y', 'issue_m'])

# drop 'zip_code' feature
data = data.drop(columns=['zip_code'])

# normalize the data
column_name = list(data.columns.values)
min_max_scaler = preprocessing.MinMaxScaler()
scaled = min_max_scaler.fit_transform(data.values)
data = pd.DataFrame(scaled)
data.columns = column_name

print (data.sample(5))
```

	loan_amnt	dti	policy_code	emp_length	purpose_0	\
701750	0.003930	4.865997e-07	0.0	0.000000	0.0	
288324	0.002501	2.105999e-07	0.0	0.000000	0.0	
726844	0.001072	2.111999e-07	0.0	0.000000	0.0	
18120	0.003215	3.799998e-08	1.0	0.111111	0.0	
598842	0.010593	3.499998e-07	0.0	0.000000	0.0	

	purpose_1	purpose_2	purpose_3	purpose_4	purpose_5	...
701750	1.0	0.0	0.0	0.0	0.0	...
288324	0.0	1.0	0.0	0.0	0.0	...
726844	0.0	0.0	0.0	1.0	0.0	...
18120	0.0	0.0	0.0	1.0	0.0	...
598842	1.0	0.0	0.0	0.0	0.0	...

	issue_m_3	issue_m_4	issue_m_5	issue_m_6	issue_m_7	issue_m_8
701750	0.0	0.0	0.0	0.0	0.0	0.0
288324	0.0	1.0	0.0	0.0	0.0	0.0
726844	0.0	0.0	0.0	0.0	0.0	0.0
18120	1.0	0.0	0.0	0.0	0.0	0.0
598842	0.0	0.0	0.0	0.0	1.0	0.0

	issue_m_9	issue_m_10	issue_m_11	issue_m_12
701750	0.0	1.0	0.0	0.0
288324	0.0	0.0	0.0	0.0
726844	0.0	0.0	1.0	0.0
18120	0.0	0.0	0.0	0.0
598842	0.0	0.0	0.0	0.0

[5 rows x 87 columns]

Modeling

As this is a classification problem, we will use a logistics regression model to predict the.

```
In [39]: from sklearn.model_selection import train_test_split

Y = data['policy_code']
data = data.drop(columns=['policy_code'])
X = data.values

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.25, random_state=10)
```

```
In [51]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report

# fit training data into a logistic classification model
logistic_predict_model = LogisticRegression()
logistic_predict_model.fit(X_train, Y_train)

# predict on the test set data
Y_predicted = logistic_predict_model.predict(X_test)
print ("The accuracy score is {:.3f}".format(accuracy_score(Y_test, Y_predicted)))
print ("\n")
print ("Classification report:")
print(classification_report(Y_test, Y_predicted))
```

The accuracy score is 0.952

Classification report:

	precision	recall	f1-score	support
0.0	0.96	0.99	0.97	186823
1.0	0.60	0.24	0.34	10364
avg / total	0.94	0.95	0.94	197187

```
In [52]: print('Coefficients:', logistic_predict_model.coef_)
```

```
Coefficients: [[-2.58051647e+01 -3.54684879e-01  3.09539972e+00  3.2724
6782e-01
 1.14629195e+00  9.84096151e-01  1.65901588e+00 -1.76462904e+00
 7.72351852e-01  2.88352646e-01 -4.47197072e-01  8.80353344e-01
 4.69619091e-01  1.00298267e+00  2.33867661e-01 -4.65560058e+00
-4.47003530e+00 -1.53562231e-01 -4.12771229e-01 -5.67588253e-01
 1.92936011e-01  3.63120395e-01  1.94059221e-01  2.46126351e-01
 1.16385432e+00 -1.30317199e-01  9.10420443e-02  3.11294771e-02
-9.98486858e-02 -3.65153215e-01  4.65879725e-01 -3.87831414e-03
-5.62932220e-01  5.42397008e-02 -5.24328787e-01 -2.52061395e-01
 4.23295145e-01  1.72678318e-01 -7.81866387e-01 -3.41373437e-01
 2.12479710e-01 -7.66000833e-02 -1.39390539e-01 -1.85593633e-01
-8.98732936e-03 -6.31446986e-01 -6.52292513e-01 -1.65434666e-01
 2.96391522e-01 -2.97814470e-01  1.77753169e-01  2.55873211e-01
-2.10861142e-01 -3.30239462e-01  2.93625517e-01 -1.59577206e-01
 4.49882534e-02 -2.26050111e-01 -1.88028899e-01  6.42011192e-02
-6.51706211e-02  3.04513569e-02  1.83800988e-01 -6.18882104e-01
 2.37913079e-01 -2.16026615e-01 -3.94116171e-01 -6.92869542e-03
 1.21652439e+00  1.05640760e+00  9.95355022e-01  2.61158352e-01
 4.64244236e-01 -7.56697356e+00 -4.59246339e-01 -4.31691302e-01
-5.59164453e-01 -4.47478179e-01  1.21997589e-01 -2.39644519e-01
-2.50420161e-01 -2.19795049e-01 -2.20250913e-01 -2.64192117e-01
-3.51086354e-01 -2.52312166e-01]]
```